

Prediction Model using Random Forest

Harry Depina

*Load the necessary libraries required for this project.
The randomForest package is used to build and train the random forest model, which is an ensemble learning method known for its robust predictive capabilities.
The caret package is utilized for splitting the dataset into training and testing sets, providing tools for data partitioning, preprocessing, and model tuning.
Additionally, the pROC package is employed to evaluate the model's performance*

```
install.packages("randomForest")
install.packages("caret")
install.packages("pROC")

library(randomForest)
library(caret)
library(pROC)

## randomForest 4.7-1.2

## Type rfNews() to see new features/changes/bug fixes.

library(caret)

## Loading required package: ggplot2

##
## Attaching package: 'ggplot2'

## The following object is masked from 'package:randomForest':
##
##     margin

## Loading required package: lattice

library(pROC)

## Type 'citation("pROC")' for a citation.

##
## Attaching package: 'pROC'

## The following objects are masked from 'package:stats':
##
##     cov, smooth, var

# Load the dataset. This dataset contains information about user interactions
# with ads for a smartphone app.
# The goal is to predict if a user will download the app after clicking the
# ad.
data <- read.csv("data.csv")
```

Examine the structure of the dataset. This helps us understand the types and formats of the variables we're working with.

The data has 8 columns: IP address, app ID, device type, OS version, channel ID, click timestamp, download timestamp, and the target variable 'downloaded'.

str(data)

```
## 'data.frame':    99999 obs. of  8 variables:
## $ X87540      : int  105560 101424 94584 68413 93663 17059 121505 192967
143636 73839 ...
## $ X12        : int   25 12 13 12 3 1 9 2 3 3 ...
## $ X1         : int    1 1 1 1 1 1 1 2 1 1 ...
## $ X13        : int   17 19 13 1 17 17 25 22 19 22 ...
## $ X497       : int   259 212 477 178 115 135 442 364 135 489 ...
## $ X11.7.17.9.30: chr  "11/7/17 13:40" "11/7/17 18:05" "11/7/17 4:58"
"11/9/17 9:00" ...
## $ X         : chr   "" "" "" "" ...
## $ X0        : int    0 0 0 0 0 0 0 0 0 0 ...
```

head(data)

```
##   X87540 X12 X1 X13 X497 X11.7.17.9.30 X X0
## 1 105560 25  1  17  259 11/7/17 13:40   0
## 2 101424 12  1  19  212 11/7/17 18:05   0
## 3  94584 13  1  13  477 11/7/17  4:58   0
## 4  68413 12  1   1  178 11/9/17  9:00   0
## 5  93663  3  1  17  115 11/9/17  1:22   0
## 6  17059  1  1  17  135 11/9/17  1:17   0
```

summary(data)

```
##      X87540          X12          X1          X13
## Min.   :      9    Min.   :  1.00    Min.   :   0.00    Min.   :   0.00
## 1st Qu.: 40552    1st Qu.:  3.00    1st Qu.:   1.00    1st Qu.: 13.00
## Median : 79827    Median : 12.00    Median :   1.00    Median : 18.00
## Mean   : 91256    Mean   : 12.05    Mean   : 21.77    Mean   : 22.82
## 3rd Qu.:118252    3rd Qu.: 15.00    3rd Qu.:   1.00    3rd Qu.: 19.00
## Max.   :364757    Max.   :551.00    Max.   :3867.00    Max.   :866.00
##      X497      X11.7.17.9.30      X      X0
## Min.   :   3.0    Length:99999    Length:99999    Min.   :0.00000
## 1st Qu.:145.0    Class :character    Class :character    1st Qu.:0.00000
## Median :258.0    Mode  :character    Mode  :character    Median :0.00000
## Mean   :268.8                                     Mean   :0.00227
## 3rd Qu.:379.0                                     3rd Qu.:0.00000
## Max.   :498.0                                     Max.   :1.00000
```

Rename columns for easier reference throughout the analysis.

This makes the code more readable and easier to interpret.

```

colnames(data) <- c("ip_address", "app_id", "device_type", "os_version",
"channel_id", "click_timestamp", "download_timestamp",
"downloaded")

# Convert the target variable 'downloaded' into a factor. This is necessary
because we are performing classification.
data$downloaded <- as.factor(data$downloaded)

# Verify the changes made to the data structure to ensure everything is set
up correctly.
str(data)

## 'data.frame': 99999 obs. of 8 variables:
## $ ip_address : int 105560 101424 94584 68413 93663 17059 121505
192967 143636 73839 ...
## $ app_id : int 25 12 13 12 3 1 9 2 3 3 ...
## $ device_type : int 1 1 1 1 1 1 1 2 1 1 ...
## $ os_version : int 17 19 13 1 17 17 25 22 19 22 ...
## $ channel_id : int 259 212 477 178 115 135 442 364 135 489 ...
## $ click_timestamp : chr "11/7/17 13:40" "11/7/17 18:05" "11/7/17 4:58"
"11/9/17 9:00" ...
## $ download_timestamp: chr "" "" "" "" ...
## $ downloaded : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...

# Set a seed for reproducibility of results. This ensures that the random
split we perform next will be the same each time we run the code.
set.seed(999)

# Split the dataset into training (70%) and testing (30%) sets using caret's
createDataPartition function.
# This prepares the data for building and evaluating the model.
trainIndex <- createDataPartition(data$downloaded, p = 0.7, list = FALSE)
train_data <- data[trainIndex, ]
test_data <- data[-trainIndex, ]

# Train the Random Forest model. We are using ip_address, app_id,
device_type, os_version, and channel_id as predictors to forecast if the app
will be downloaded.
# The model is set to measure variable importance.
rf_model <- randomForest(downloaded ~ ip_address + app_id + device_type +
os_version + channel_id,
data = train_data, importance = TRUE)

# Print the model summary. This includes details like the number of trees,
the variables tried at each split, and the Out-Of-Bag
(OOB) error estimate.
# The OOB error rate provides an estimate of the model's error on unseen
data.
print(rf_model)

```

```
##
## Call:
## randomForest(formula = downloaded ~ ip_address + app_id + device_type +
+ os_version + channel_id, data = train_data, importance = TRUE)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 2
##
##           OOB estimate of error rate: 0.2%
## Confusion matrix:
##           0  1  class.error
## 0 69823 18 0.0002577283
## 1   122 37 0.7672955975

# Predict on the test data using the trained Random Forest model. This step
# gives us the predicted download statuses for the test data.
predictions <- predict(rf_model, test_data)

# Display the first few predictions to get a sense of the model's output.
head(predictions)

##  2  9 15 22 23 30
##  0  0  0  0  0  0
## Levels: 0 1

# Evaluate the model's performance using a confusion matrix. This matrix
# compares the predicted values to the actual values in the test data.
# Metrics like accuracy, sensitivity, and specificity are calculated to
# assess how well the model performed.
confusion_matrix <- confusionMatrix(predictions, test_data$downloaded)
print(confusion_matrix)

## Confusion Matrix and Statistics
##
##           Reference
## Prediction      0      1
##           0 29927     54
##           1      4     14
##
##           Accuracy : 0.9981
##           95% CI : (0.9975, 0.9985)
##           No Information Rate : 0.9977
##           P-Value [Acc > NIR] : 0.1228
##
##           Kappa : 0.3249
##
##           Mcnemar's Test P-Value : 1.243e-10
##
##           Sensitivity : 0.9999
##           Specificity : 0.2059
##           Pos Pred Value : 0.9982
```

```
##          Neg Pred Value : 0.7778
##          Prevalence : 0.9977
##          Detection Rate : 0.9976
##    Detection Prevalence : 0.9994
##          Balanced Accuracy : 0.6029
##
##          'Positive' Class : 0
##
```

The model achieves an accuracy of around 99.81%. However, the sensitivity and specificity metrics reveal an imbalance, as the model is better at identifying non-downloads (class "0") than downloads (class "1").
This imbalance is common in classification problems with skewed class distributions.

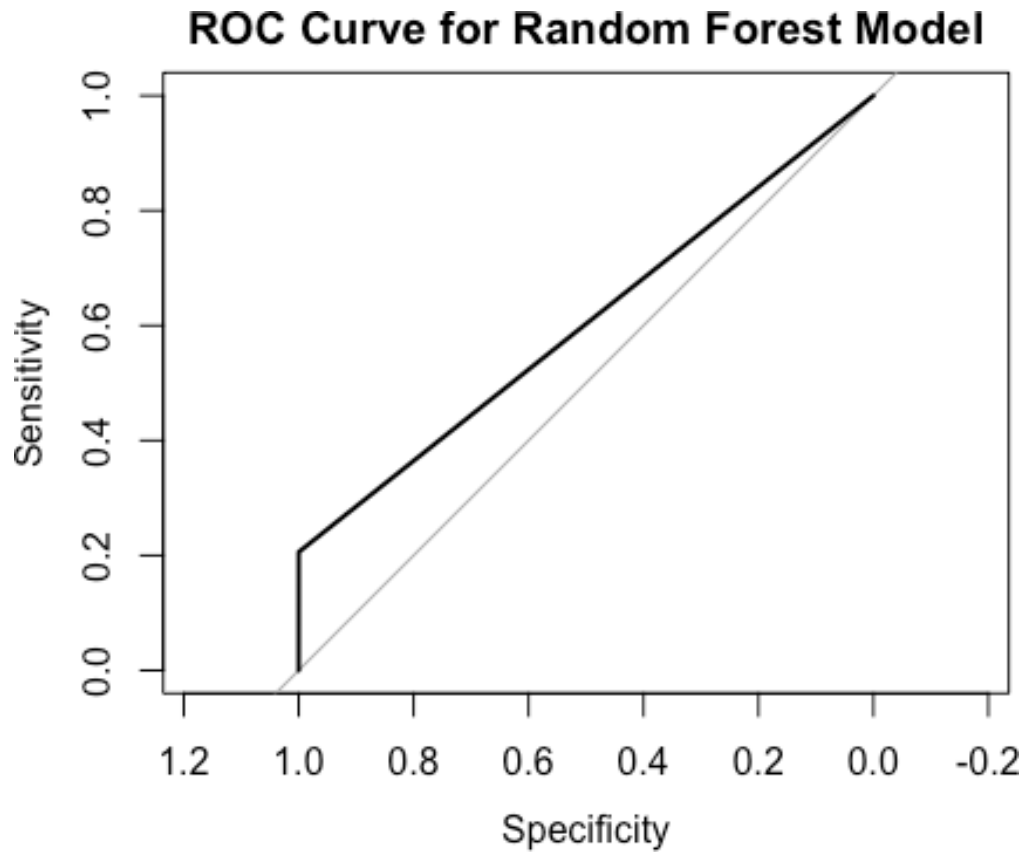
Plot the ROC curve to visualize the trade-off between the true positive rate and false positive rate.
A higher Area Under the Curve (AUC) value indicates better model performance.

```
roc_curve <- roc(test_data$downloaded, as.numeric(predictions))
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
plot(roc_curve, main = "ROC Curve for Random Forest Model")
```



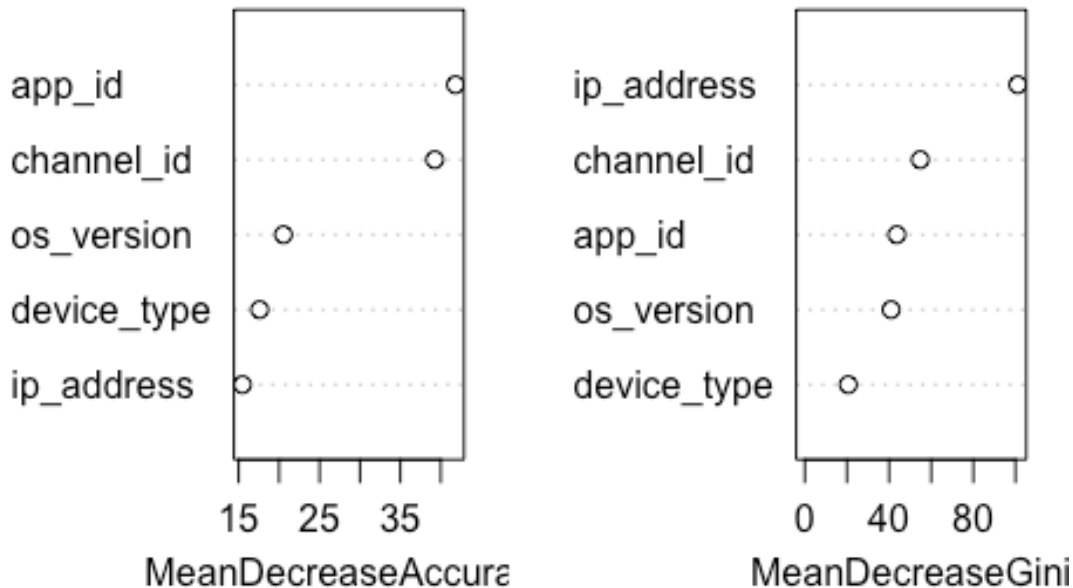
The ROC curve shows that the model has a relatively good ability to differentiate between downloads and non-downloads, as it lies above the 45-degree line, indicating predictive power.

Plot feature importance to identify which variables had the most impact on the predictions.

This helps us understand the driving factors behind user app downloads.

```
varImpPlot(rf_model, main = "Feature Importance in Random Forest Model")
```

Feature Importance in Random Forest Model



```
# Extract and display the feature importance values from the model.
# The MeanDecreaseAccuracy metric shows the importance of each variable in
# maintaining the model's accuracy.
# The MeanDecreaseGini metric indicates each variable's role in maintaining
# node purity across the trees in the forest.
importance_values <- importance(rf_model)
print(importance_values)
```

	0	1	MeanDecreaseAccuracy	MeanDecreaseGini
ip_address	-6.585658	44.14471	15.46904	101.12717
app_id	36.632892	63.17607	41.82380	43.70156
device_type	15.713511	12.80190	17.59922	20.66655
os_version	16.456376	24.58729	20.54581	40.98300
channel_id	35.964003	32.57665	39.25806	54.81954

```
# Key findings from feature importance analysis:
# The 'app_id' and 'channel_id' variables are among the top predictors,
# suggesting that certain apps and ad channels are more effective in driving
# downloads.
```

```
# This insight can help marketers optimize ad targeting based on the
# performance of specific channels and app IDs.
```

```
# Conclusion
```

In this project, we successfully developed a Random Forest model to predict whether a user will download an app after clicking on an ad.

The model achieved high accuracy but displayed a bias towards non-download predictions, as seen from the specificity metric.

The feature importance analysis indicates that app_id and channel_id are key factors in predicting downloads,

offering actionable insights for improving ad targeting strategies.

This concludes the analysis. The model's predictions, accuracy metrics, ROC curve, and feature importance collectively provide a comprehensive view

of the factors influencing app downloads.